

Code and output screenshots from python notebook

implementation of LSTM Model to predict the soil temperature for the next 6 months, using Python in jupyter notebook

Student Name: Gargi Darade

Subject/Class: CMPS 451 Artificial Intelligence

#Step 1:- Import the required libraries

#Numpy for statistical computations

#Matplotlib to plot the graph

#make_blobs from sklearn.datasets

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

#from sklearn.datasets import make_blobs

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import MinMaxScaler

#!pip install tensorflow

import tensorflow as tf

data_path = 'psspredit_new01.csv'

create pandas dataframe

df = pd.read_csv(data_path)

remove spaces on the column

df.columns = df.columns.str.lstrip()

df.columns = df.columns.str.rstrip()

print out sample dataset

print(len(df))

df.head()

Out put

date	naturaltemperature_5	windspeedscalar	temperature_1p5	
	temperature_2	relativehumidity	stationpressure	solarradiation
	windspeed_20f			
0	2001/1/2 2:00	0.939 -4.308	4.611 4.599 88.60	4.156 0.335 5.325
1	2001/1/2 3:00	0.950 -4.554	4.283 4.273 52.06	3.964 0.494 5.292
2	2001/1/2 4:00	0.958 -4.513	3.359 3.347 70.40	4.891 0.230 3.822
3	2001/1/2 5:00	0.969 -4.670	3.620 3.606 64.33	4.952 0.276 4.345
4	2001/1/2 6:00	0.976 -5.170	3.819 3.807 62.07	4.699 0.327 4.508

check number of nan values in dataframe

df.isna().sum()

Out put

```
date                0
naturaltemperature_5  0
windspeedscalar     0
temperature_1p5     0
temperature_2       0
relativehumidity    0
stationpressure     0
solarradiation      0
windspeed_20f       0
dtype: int64
```

plot to see the soil temperature variance

#conda update --all

#pip install -U seaborn

plt.figure(figsize=(5, 5))

sns.histplot(df['naturaltemperature_5'],bins=[i for i in range(0,61,5)], kde=False)

plt.title("Distribution of Soil Temperatures")

plt.grid()

#plt.tight_layout()

plt.savefig("SoilTemp2.png")

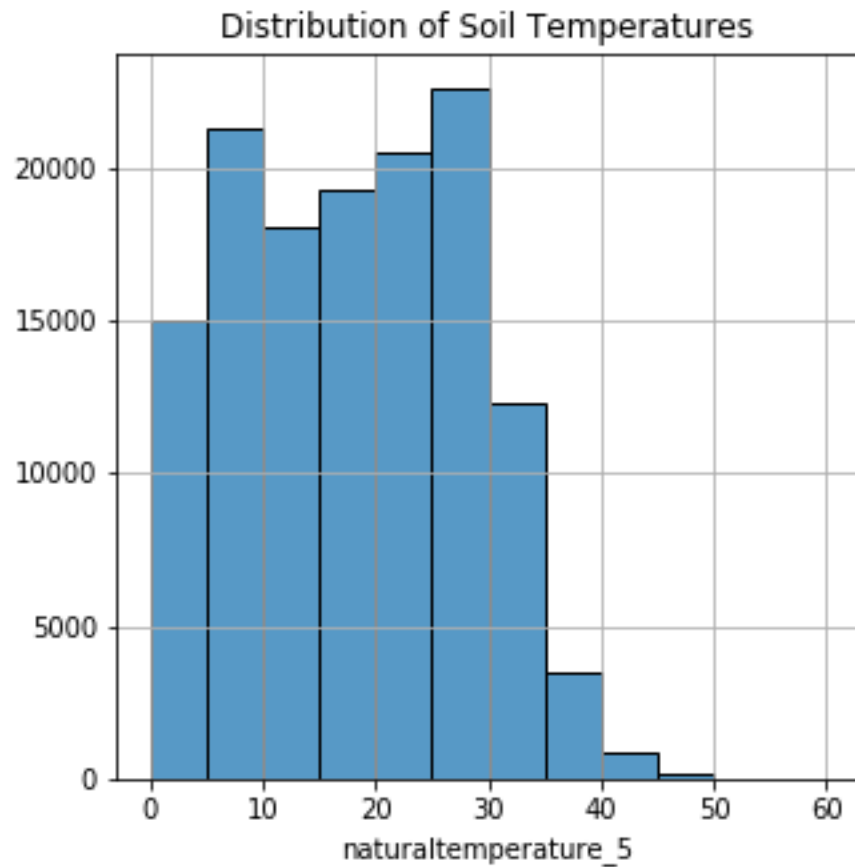
plt.show()

#sns.histplot(df['windspeedscalar'],bins=[i for i in range(0,61,5)], kde=False)

#plt.title("Distribution of Wind Speed")

#plt.grid()

#plt.show()



```
f, axes = plt.subplots(2, 4, figsize=(7, 7), sharex=True)
```

```
sns.despine(left=True)
```

```
# Plot a simple distribution of the desired columns
```

```
sns.distplot(df['naturaltemperature_5'], color="b", ax=axes[0, 0])
```

```
sns.distplot(df['windspeedscalar'], color="m", ax=axes[0, 1])
```

```
sns.distplot(df['temperature_1p5'], color="r", ax=axes[0, 2])
```

```
sns.distplot(df['temperature_2'], color="g", ax=axes[0, 3])
```

```
sns.distplot(df['relativehumidity'], color="b", ax=axes[1, 0])
```

```
sns.distplot(df['stationpressure'], color="m", ax=axes[1, 1])
```

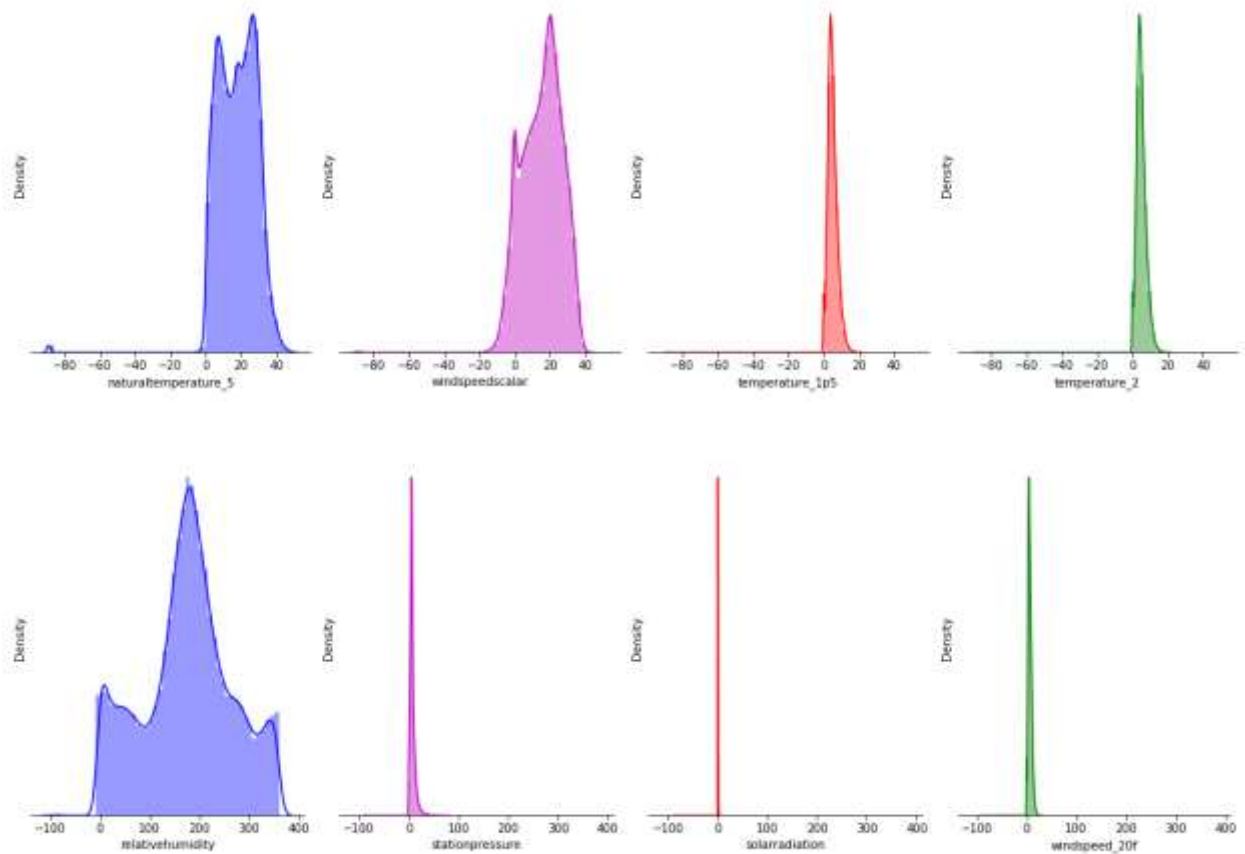
```
sns.distplot(df['solarradiation'], color="r", ax=axes[1, 2])
```

```
sns.distplot(df['windspeed_20f'], color="g", ax=axes[1, 3])
```

```
plt.setp(axes, yticks=[])
```

```
plt.tight_layout()
```

```
#plt.grid()
#plt.show()
plt.savefig("sample.png")
```



```
# set data index as datetime column
df.index = pd.to_datetime(df['date'])
# filter the columns by only the required_columns
required_cols = ['naturaltemperature_5', 'windspeedscalar', 'temperature_1p5', 'temperature_2',
'relativehumidity', 'stationpressure', 'solarradiation', 'windspeed_20f']
df = df[required_cols]
df.head()
```

Out put

naturaltemperature_5	windspeedscalar	temperature_1p5	temperature_2
relativehumidity	stationpressure	solarradiation	windspeed_20f
date			
2001-01-02 02:00:00	0.939 -4.308	4.611 4.599 88.60	4.156 0.335 5.325
2001-01-02 03:00:00	0.950 -4.554	4.283 4.273 52.06	3.964 0.494 5.292
2001-01-02 04:00:00	0.958 -4.513	3.359 3.347 70.40	4.891 0.230 3.822

```
2001-01-02 05:00:00    0.969 -4.670    3.620 3.606 64.33 4.952 0.276 4.345
2001-01-02 06:00:00    0.976 -5.170    3.819 3.807 62.07 4.699 0.327 4.508
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
# Normalize the data
```

```
scaler = MinMaxScaler()
```

```
scaled_data = scaler.fit_transform(df)
```

```
# Define sequence length and features
```

```
sequence_length = 24 # Number of time steps in each sequence
```

```
num_features = len(df.columns)
```

```
# Create sequences and corresponding labels
```

```
sequences = []
```

```
labels = []
```

```
for i in range(len(scaled_data) - sequence_length):
```

```
    seq = scaled_data[i:i+sequence_length]
```

```
    label = scaled_data[i+sequence_length][0] # 'naturaltemperature_5' column index
```

```
    sequences.append(seq)
```

```
    labels.append(label)
```

```
# Convert to numpy arrays
```

```
sequences = np.array(sequences)
```

```
labels = np.array(labels)
```

```
# Split into train and test sets
```

```
train_size = int(0.8 * len(sequences))
```

```
train_x, test_x = sequences[:train_size], sequences[train_size:]
```

```
train_y, test_y = labels[:train_size], labels[train_size:]
```

```
print("Train X shape:", train_x.shape)
```

```
print("Train Y shape:", train_y.shape)
```

```
print("Test X shape:", test_x.shape)
```

```
print("Test Y shape:", test_y.shape)
```

```
Train X shape: (105590, 10, 8)
```

```
Train Y shape: (105590,)
```

```
Test X shape: (26398, 10, 8)
```

```
Test Y shape: (26398,)
```

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

# Create the LSTM model
model = Sequential()

# Add LSTM layers with dropout
model.add(LSTM(units=128, input_shape=(train_x.shape[1], train_x.shape[2]),
return_sequences=True))
model.add(Dropout(0.2))

model.add(LSTM(units=64, return_sequences=True))
model.add(Dropout(0.2))

model.add(LSTM(units=32, return_sequences=False))
model.add(Dropout(0.2))

# Add a dense output layer
model.add(Dense(units=1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 24, 128)	70144
dropout_3 (Dropout)	(None, 24, 128)	0
lstm_4 (LSTM)	(None, 24, 64)	49408
dropout_4 (Dropout)	(None, 24, 64)	0
lstm_5 (LSTM)	(None, 32)	12416
dropout_5 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 1)	33

Total params: 132,001

Trainable params: 132,001
Non-trainable params: 0

Define callbacks

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
model_checkpoint =
ModelCheckpoint('C:/Users/Manomay/Downloads/Gargi/CMPS451AI/assignmentsAI/soilTempLSTM/
best_model_weights.h5', monitor='val_loss', save_best_only=True)
```

Train the model

```
history = model.fit(
    train_x, train_y,
    epochs=100,
    batch_size=64,
    validation_split=0.2, # Use part of the training data as validation
    callbacks=[early_stopping, model_checkpoint]
)
```

```
Epoch 1/100
1320/1320 [=====] - 23s 14ms/step - loss: 0.0075
- val_loss: 8.7919e-04
Epoch 2/100
1320/1320 [=====] - 19s 15ms/step - loss: 0.0027
- val_loss: 8.0958e-04
Epoch 3/100
1320/1320 [=====] - 21s 16ms/step - loss: 0.0011
- val_loss: 4.3629e-04
Epoch 4/100
1320/1320 [=====] - 19s 15ms/step - loss: 4.1244e
-04 - val_loss: 6.4938e-04
Epoch 5/100
1320/1320 [=====] - 20s 15ms/step - loss: 2.3890e
-04 - val_loss: 6.9732e-04
Epoch 6/100
1320/1320 [=====] - 19s 14ms/step - loss: 1.8967e
-04 - val_loss: 2.5110e-04
Epoch 7/100
1320/1320 [=====] - 19s 14ms/step - loss: 1.6482e
-04 - val_loss: 3.1614e-04
Epoch 8/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.6010e
-04 - val_loss: 2.4190e-04
Epoch 9/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.4313e
-04 - val_loss: 2.2215e-04
Epoch 10/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.4497e
-04 - val_loss: 2.2920e-04
```

```

Epoch 11/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.4020e
-04 - val_loss: 3.1838e-04
Epoch 12/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.2007e
-04 - val_loss: 3.1465e-04
Epoch 13/100
1320/1320 [=====] - 19s 15ms/step - loss: 1.2765e
-04 - val_loss: 3.8833e-04
Epoch 14/100
1320/1320 [=====] - 20s 15ms/step - loss: 1.3454e
-04 - val_loss: 5.5418e-04
Epoch 11/100
1342/1342 [=====] - 40s 30ms/step - loss: 1.3989e
-04 - val_loss: 2.6575e-04
Epoch 12/100
1342/1342 [=====] - 40s 30ms/step - loss: 1.2990e
-04 - val_loss: 2.8687e-04
Epoch 13/100
1342/1342 [=====] - 41s 30ms/step - loss: 1.2928e
-04 - val_loss: 2.4973e-04
Epoch 14/100
1342/1342 [=====] - 41s 30ms/step - loss: 1.2714e
-04 - val_loss: 2.2059e-04
Epoch 15/100
1342/1342 [=====] - 41s 30ms/step - loss: 1.2743e
-04 - val_loss: 3.3258e-04
Epoch 16/100
1342/1342 [=====] - 44s 33ms/step - loss: 1.2538e
-04 - val_loss: 4.1259e-04
Epoch 17/100
1342/1342 [=====] - 43s 32ms/step - loss: 1.1183e
-04 - val_loss: 5.1869e-04
Epoch 18/100
1342/1342 [=====] - 42s 32ms/step - loss: 1.0847e
-04 - val_loss: 5.0255e-04
Epoch 19/100
1342/1342 [=====] - 43s 32ms/step - loss: 1.1605e
-04 - val_loss: 5.3437e-04

```

Evaluate the best model on the test set

```
best_model = tf.keras.models.load_model('best_model_weights.h5')
```

```
test_loss = best_model.evaluate(test_x, test_y)
```

#provides an estimate of how well the trained model generalizes to new, unseen data

```
print("Test Loss:", test_loss)
```

```

825/825 [=====] - 4s 4ms/step - loss: 1.7681e-04
Test Loss: 0.00017681172175798565

```

Plot training & validation loss values


```
plt.plot(history.history['loss'])
```



```

plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.savefig("MLoss.png")
plt.show()

```



```

from sklearn.metrics import mean_absolute_error, mean_squared_error

```

```

# Assuming you have trained the model and have the 'best_model' object
# Also, 'test_x' and 'test_y' should be available

```

```

# Predict temperatures using the trained model
predictions = best_model.predict(test_x)

```

```

# Calculate evaluation metrics
mae = mean_absolute_error(test_y, predictions)
mse = mean_squared_error(test_y, predictions)
rmse = np.sqrt(mse)

```

```

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)

```

```
print("Root Mean Squared Error (RMSE):", rmse)
```

```
825/825 [=====] - 4s 4ms/step  
Mean Absolute Error (MAE): 0.004908275340278715  
Mean Squared Error (MSE): 0.00017681170242666888  
Root Mean Squared Error (RMSE): 0.013297056156408037
```

```
# y_true values
```

```
test_y_copies = np.repeat(test_y.reshape(-1, 1), test_x.shape[-1], axis=-1)
```

```
true_temp = scaler.inverse_transform(test_y_copies)[: ,0] # 'naturaltemperature_5' column index
```

```
# predicted values
```

```
prediction = best_model.predict(test_x)
```

```
prediction_copies = np.repeat(prediction, num_features, axis=-1)
```

```
predicted_temp = scaler.inverse_transform(prediction_copies)[: ,0] # 'naturaltemperature_5' column  
index
```

```
825/825 [=====] - 3s 4ms/step
```

```
# Plotting predicted and actual temperatures
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(df.index[-100:], true_temp[-100:], label='Actual')
```

```
plt.plot(df.index[-100:], predicted_temp[-100:], label='Predicted')
```

```
plt.title('Soil Temperature Prediction vs Actual')
```

```
plt.xlabel('Time')
```

```
plt.ylabel('Temperature')
```

```
plt.legend()
```

```
plt.savefig("STempPred.png")
```

```
plt.show()
```

